

student-performance-prediction-22

September 4, 2025

```
[1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

import warnings
warnings.filterwarnings("ignore")
```

```
[2]: data = pd.read_csv('/kaggle/input/
↳student-performance-multiple-linear-regression/Student_Performance.csv')
```

```
[3]: # Display Top 5 Rows of the DataFrame
data.head()
```

```
[3]:  Hours Studied  Previous Scores  Extracurricular Activities  Sleep Hours  \
0              7              99                      Yes          9
1              4              82                      No           4
2              8              51                      Yes          7
3              5              52                      Yes          5
4              7              75                      No           8
```

```
      Sample Question Papers Practiced  Performance Index
0                                1          91.0
1                                2          65.0
2                                2          45.0
3                                2          36.0
4                                5          66.0
```

```
[5]: data.shape
```

```
[5]: (10000, 6)
```

```
[6]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 10000 entries, 0 to 9999
```

```
Data columns (total 6 columns):
```

```
#      Column                                Non-Null Count  Dtype
---  -
```

```

0    Hours Studied          10000 non-null  int64
1    Previous Scores        10000 non-null  int64
2    Extracurricular Activities 10000 non-null  object
3    Sleep Hours            10000 non-null  int64
4    Sample Question Papers Practiced 10000 non-null  int64
5    Performance Index       10000 non-null  float64
dtypes: float64(1), int64(4), object(1)
memory usage: 468.9+ KB

```

```
[7]: data.isnull().sum()
```

```

[7]: Hours Studied          0
     Previous Scores        0
     Extracurricular Activities 0
     Sleep Hours            0
     Sample Question Papers Practiced 0
     Performance Index       0
     dtype: int64

```

```
[10]: # Check For Duplicates
      data.duplicated().sum()
```

```
[10]: 127
```

```
[11]: # Drop duplicates
      data.drop_duplicates(keep='first', inplace=True)
```

```
[12]: data.shape
```

```
[12]: (9873, 6)
```

```
[13]: # Quick Statistics
      data.describe()
```

```

[13]:      Hours Studied  Previous Scores  Sleep Hours  \
count      9873.000000      9873.000000  9873.000000
mean         4.992100        69.441102     6.531652
std          2.589081        17.325601     1.697683
min           1.000000        40.000000     4.000000
25%           3.000000        54.000000     5.000000
50%           5.000000        69.000000     7.000000
75%           7.000000        85.000000     8.000000
max           9.000000        99.000000     9.000000

      Sample Question Papers Practiced  Performance Index
count      9873.000000      9873.000000
mean         4.583004        55.216651
std          2.867202        19.208570

```

min	0.000000	10.000000
25%	2.000000	40.000000
50%	5.000000	55.000000
75%	7.000000	70.000000
max	9.000000	100.000000

```
[14]: data.head()
```

```
[14]:
```

	Hours Studied	Previous Scores	Extracurricular Activities	Sleep Hours	\
0	7	99	Yes	9	
1	4	82	No	4	
2	8	51	Yes	7	
3	5	52	Yes	5	
4	7	75	No	8	

	Sample Question Papers Practiced	Performance Index
0	1	91.0
1	2	65.0
2	2	45.0
3	2	36.0
4	5	66.0

```
[15]: data['Extracurricular Activities'].value_counts()
```

```
[15]: Extracurricular Activities
No      4986
Yes     4887
Name: count, dtype: int64
```

```
[16]: data['Extracurricular Activities'] = data['Extracurricular Activities'].
      ↪map({'Yes':1, 'No':0})
```

```
[17]: data.head()
```

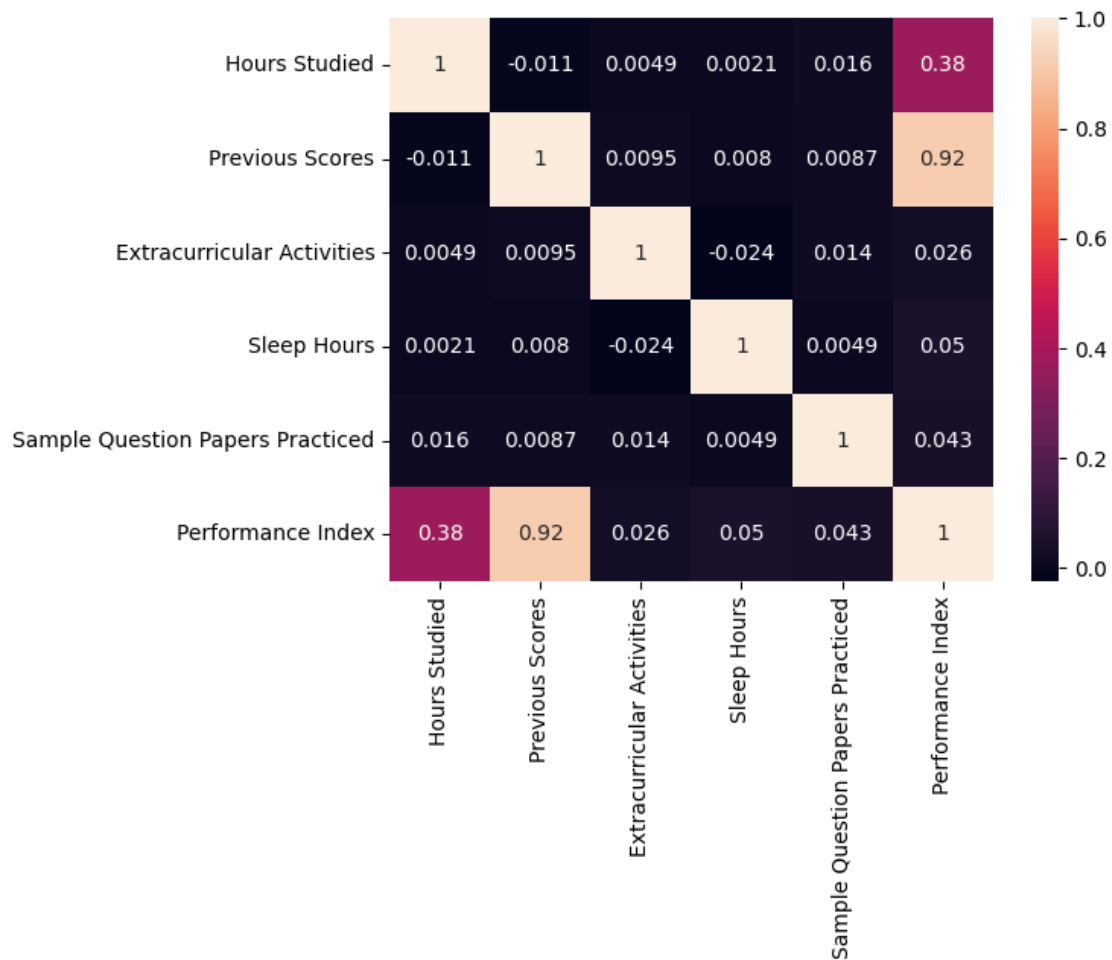
```
[17]:
```

	Hours Studied	Previous Scores	Extracurricular Activities	Sleep Hours	\
0	7	99	1	9	
1	4	82	0	4	
2	8	51	1	7	
3	5	52	1	5	
4	7	75	0	8	

	Sample Question Papers Practiced	Performance Index
0	1	91.0
1	2	65.0
2	2	45.0
3	2	36.0
4	5	66.0

```
[18]: sns.heatmap(data.corr(), annot=True )
```

```
[18]: <Axes: >
```



```
[19]: data.corr()
```

```
[19]:
```

	Hours Studied	Previous Scores	\
Hours Studied	1.000000	-0.010676	
Previous Scores	-0.010676	1.000000	
Extracurricular Activities	0.004899	0.009534	
Sleep Hours	0.002131	0.007975	
Sample Question Papers Practiced	0.015740	0.008719	
Performance Index	0.375332	0.915135	

	Extracurricular Activities	Sleep Hours	\
Hours Studied	0.004899	0.002131	
Previous Scores	0.009534	0.007975	

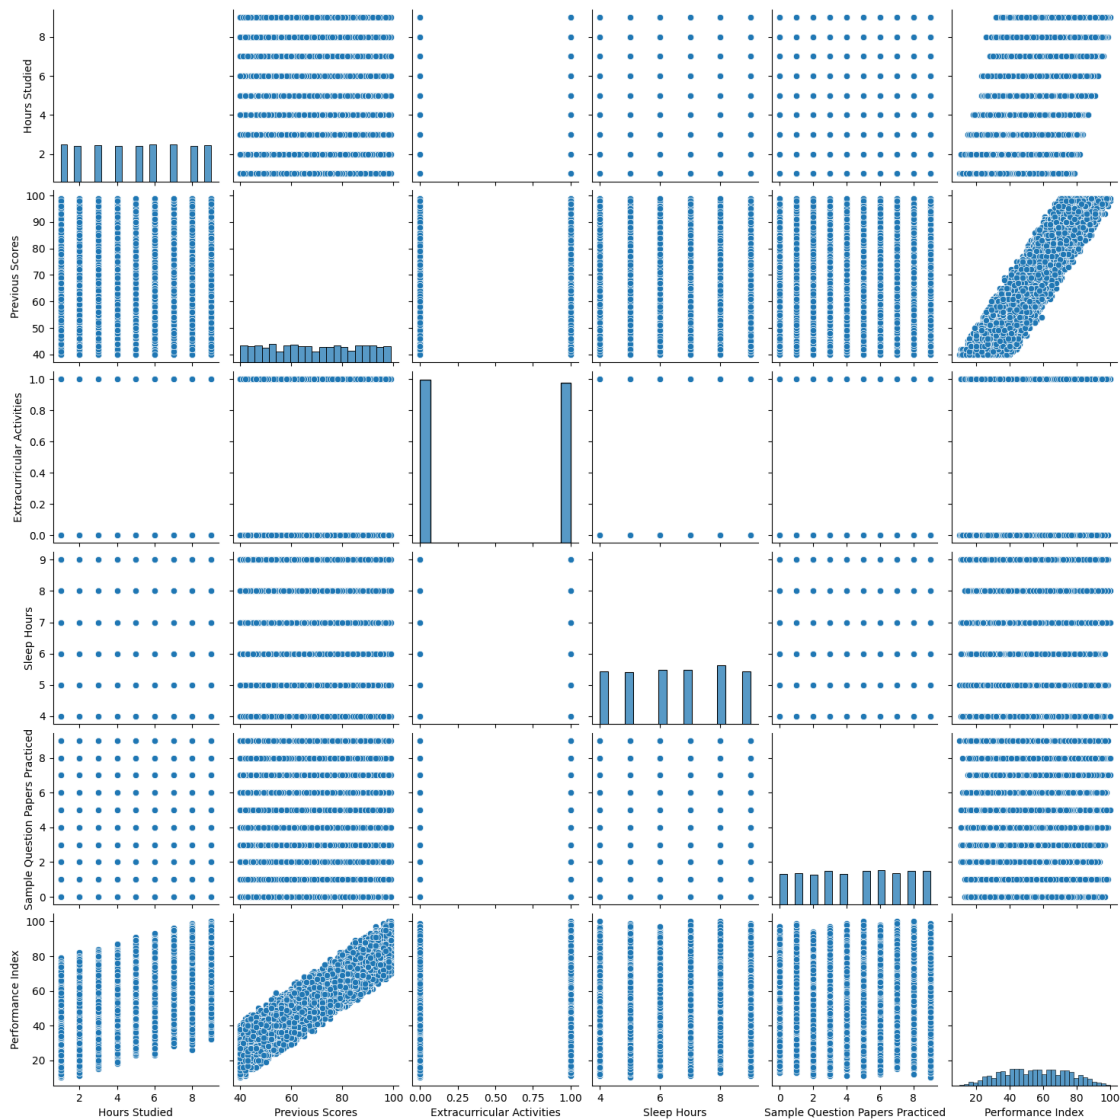
Extracurricular Activities	1.000000	-0.024008
Sleep Hours	-0.024008	1.000000
Sample Question Papers Practiced	0.013839	0.004907
Performance Index	0.026075	0.050352

	Sample Question Papers Practiced \
Hours Studied	0.015740
Previous Scores	0.008719
Extracurricular Activities	0.013839
Sleep Hours	0.004907
Sample Question Papers Practiced	1.000000
Performance Index	0.043436

	Performance Index
Hours Studied	0.375332
Previous Scores	0.915135
Extracurricular Activities	0.026075
Sleep Hours	0.050352
Sample Question Papers Practiced	0.043436
Performance Index	1.000000

```
[20]: sns.pairplot(data)
```

```
[20]: <seaborn.axisgrid.PairGrid at 0x7d092317ead0>
```



```
[21]: data.head()
```


```
[21]:
```

	Hours Studied	Previous Scores	Extracurricular Activities	Sleep Hours	\
0	7	99	1	9	
1	4	82	0	4	
2	8	51	1	7	
3	5	52	1	5	
4	7	75	0	8	

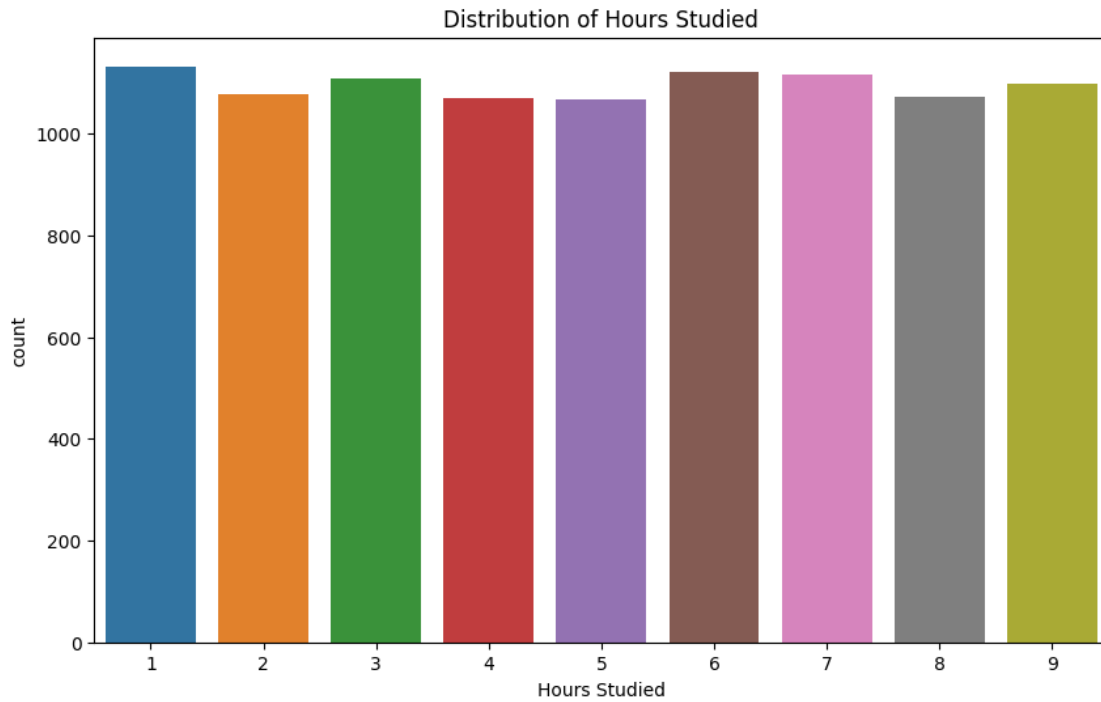
	Sample Question Papers Practiced	Performance Index
0	1	91.0
1	2	65.0
2	2	45.0

3	2	36.0
4	5	66.0

```
[23]: plt.figure(figsize=(10,6))
sns.histplot(x='Previous Scores',data=data)
plt.title('Distribution of Previous Scores')
plt.show()
```



```
[24]: plt.figure(figsize=(10,6))
sns.countplot(x='Hours Studied',data=data)
plt.title('Distribution of Hours Studied')
plt.show()
```



```
[25]: X = data.drop('Performance Index', axis=1)
      y = data['Performance Index']
```

```
[26]: from sklearn.model_selection import train_test_split
```

```
[27]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪ random_state=42)
```

```
[28]: from sklearn.linear_model import LinearRegression
```

```
[29]: lr = LinearRegression()
      lr.fit(X_train, y_train)
```

```
[29]: LinearRegression()
```

```
[30]: y_pred = lr.predict(X_test)
```

```
[31]: from sklearn.metrics import r2_score
```

```
[32]: print(r2_score(y_test, y_pred))
```

```
0.9884301209927054
```

```
[33]: before_fs_r2 = r2_score(y_test, y_pred)
```



```
[35]: before_fs_r2
```

```
[35]: 0.9884301209927054
```

```
[36]: data.head()
```

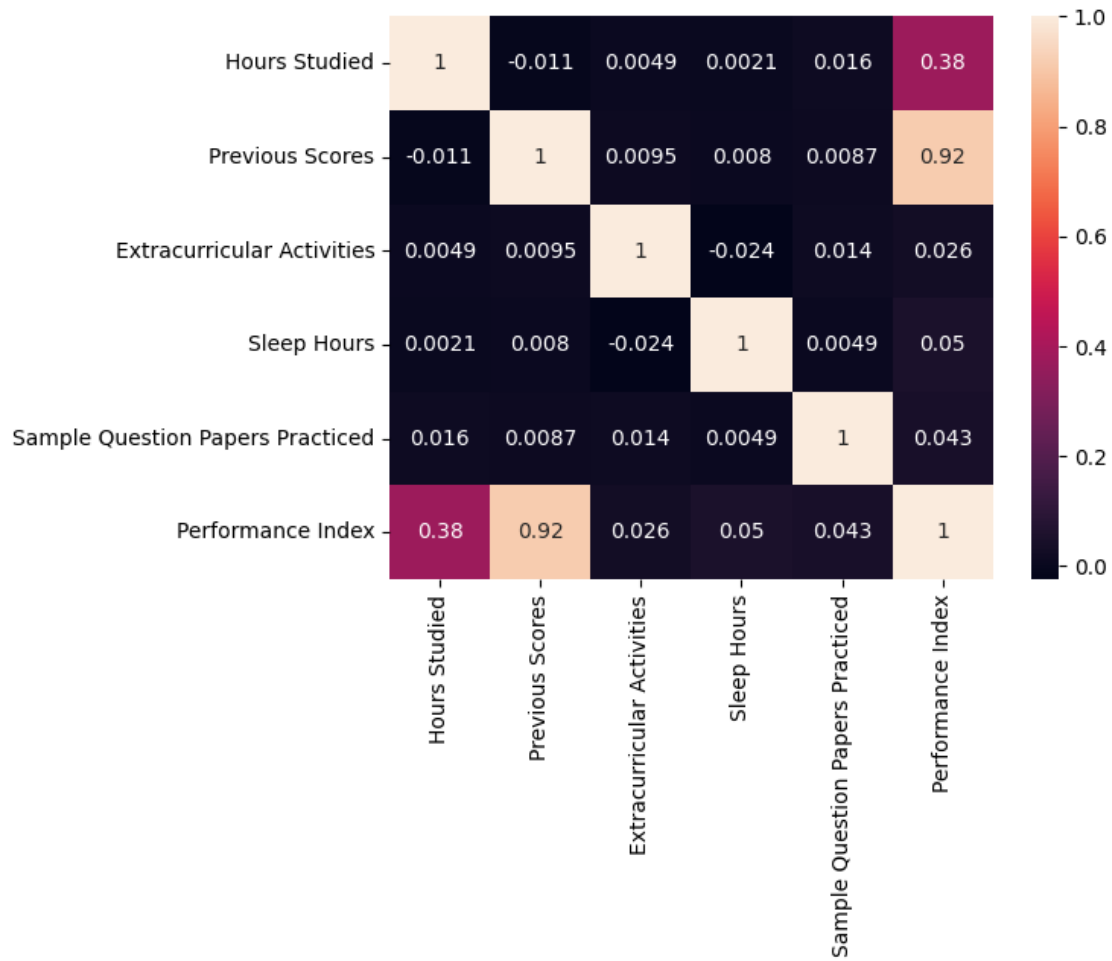
```
[36]:
```

	Hours Studied	Previous Scores	Extracurricular Activities	Sleep Hours	\
0	7	99	1	9	
1	4	82	0	4	
2	8	51	1	7	
3	5	52	1	5	
4	7	75	0	8	

	Sample Question Papers Practiced	Performance Index
0	1	91.0
1	2	65.0
2	2	45.0
3	2	36.0
4	5	66.0

```
[38]: sns.heatmap(data.corr(), annot=True)
```

```
[38]: <Axes: >
```



```
[37]: df = data.drop(['Extracurricular Activities', 'Sleep Hours', 'Sample Question Papers Practiced'], axis=1)
```

```
[39]: df
```

```
[39]:
```

	Hours Studied	Previous Scores	Performance Index
0	7	99	91.0
1	4	82	65.0
2	8	51	45.0
3	5	52	36.0
4	7	75	66.0
...
9995	1	49	23.0
9996	7	64	58.0
9997	6	83	74.0
9998	9	97	95.0
9999	7	74	64.0

[9873 rows x 3 columns]

```
[40]: X_df = df.drop('Performance Index', axis=1)
      y_df = df['Performance Index']
```

```
[41]: X_train, X_test, y_train, y_test = train_test_split(X_df, y_df, test_size=0.2,
      ↪random_state=42)
```

```
[42]: lr_f = LinearRegression()
      lr_f.fit(X_train, y_train)
```

```
[42]: LinearRegression()
```

```
[43]: yf_pred = lr_f.predict(X_test)
```

```
[44]: after_fs_r2 = r2_score(y_test, yf_pred)
```

```
[45]: print( after_fs_r2)
```

0.9850233951895029

```
[46]: Result = pd.Series({'Before Feature Selection':before_fs_r2,'After Feature_
      ↪Selection':after_fs_r2 })
```

```
[47]: Result
```

```
[47]: Before Feature Selection    0.988430
      After Feature Selection    0.985023
      dtype: float64
```

```
[50]: from sklearn.preprocessing import PolynomialFeatures
      from sklearn.metrics import r2_score
```

```
degree = 2
poly = PolynomialFeatures(degree=degree)
X_train_poly = poly.fit_transform(X_train)
X_test_poly = poly.transform(X_test)
```

```
y_pred = poly_model.predict(X_test_poly)
print("R2 Score:", r2_score(y_test, y_pred))
```

R² Score: 0.9850343841670226

```
[ ]:
```